

Machine Learning

Section 5: Continuous Probabilities

第5节：连续概率

Stefan Harmeling

13. October 2021

某一不确定量 p 的先验概率（Prior probability）分布是在考虑“观测数据”前，能表达 p 不确定性的概率分布。它旨在描述这个不确定量的不确定程度，而不是这个不确定量的随机性。

在贝叶斯统计中，一个随机事件或者一个不确定事件的后验概率（Posterior probability）是在考虑和给出相关证据或数据后所得到的条件概率。同样，后验概率分布是一个未知量（视为随机变量）基于试验和调查后得到的概率分布。“后验”在本文中代表考虑了被测试事件的相关证据。

Probabilities on continuous variables

连续变量的概率

From discrete to continuous variables

from Murphy 2012, p32, from Jaynes 2003, p107

从离散变量到连续变量

- ▶ let X be real-valued variable 设 X 为实值变量
- ▶ define propositions $A = (X \leq a)$, $B = (X \leq b)$ and $W = (a < X \leq b)$ with $a < b$
- ▶ note that $E_B = E_{A \vee W}$
- ▶ note that A and W are mutually exclusive, thus sum rule:

$$p(W) = p(B) - p(A)$$

- ▶ with $F(x) := p(X \leq x)$ and $f(x) = \frac{d}{dx} F(x)$ we get

$$p(a < X \leq b) = F(b) - F(a) = \int_a^b f(x) dx$$

- ▶ F is called *cumulative distribution function* (CDF) and f is called *probability density function* (PDF)

概率密度函数(PDF)

Discrete vs. continuous probabilities (1)

概率论作为命题逻辑的延伸

Probability theory as an extension of propositional logic

- ▶ finite set of propositional variables $A, B, \dots, Z \in \{0, 1\}$ jointly ranging over all boolean assignments 样本空间 $\Omega = \{\text{所有布尔赋值}\}$ 。
- ▶ sample space $\Omega = \{\text{all boolean assignments}\}$
- ▶ probability mass function $f: \Omega \rightarrow [0, 1]$, such that $\sum_{\omega \in \Omega} f(\omega) = 1$

Discrete probability theory (includes the previous case) 离散概率

- ▶ random variable ranging in a discrete set, e.g. $\{0, 1, 2, \dots\}$
- ▶ sample space $\Omega = \{0, 1, 2, \dots\}$
- ▶ *probability mass function* $f: \Omega \rightarrow [0, 1]$, such that $\sum_{\omega \in \Omega} f(\omega) = 1$

Continuous probability theory

- ▶ random variable ranging in a continuous set, e.g. real numbers \Re
- ▶ sample space $\Omega = \Re$
- ▶ *probability density function* $f: \Omega \rightarrow \Re_+$, such that $\int_{\omega \in \Omega} f(\omega) d\omega = 1$

Discrete vs. continuous probabilities (2)

概率衡量样本空间的子集 $E \subset \Omega$ 的质量

Probabilities measure the mass of subsets $E \subset \Omega$ of the sample space.

Discrete probabilities

- ▶ e.g. $\Omega = \mathbb{N}$
- ▶ probability mass function
 $f : \mathbb{N} \rightarrow [0, 1]$, such that
 $\sum_n f(n) = 1$
- ▶ probability $p(E) = \sum_{n \in E} f(n)$
- ▶ small letter p

Continuous probabilities

- ▶ e.g. $\Omega = \mathfrak{R}$
- ▶ probability density function
(PDF) $f : \mathfrak{R} \rightarrow \mathfrak{R}_+$, such that
 $\int f(x) dx = 1$
- ▶ probability $P(E) = \int_E f(x) dx$
- ▶ large letter P

小字母p

Rules for continuous variables

Theorem 5.1 (rules for PDFs)

The standard rules of probability theory do hold for PDFs,

$$f(x, y) = f(x|y) f(y) \quad \text{product rule}$$

$$f(x) = \int f(x, y) dy \quad \text{sum rule}$$

*For that reason we often write p for PDFs. Sums turn into integrals.
Note that the product rule implies Bayes' rule.*

由于这个原因，我们经常把 p 写成PDF。和变成了积分。请注意，乘积规则意味着贝叶斯规则。

Different way to think about the probability rules

- ▶ consider three random variables A, B, C with values a, b, c
- ▶ for probabilistic inference we need the joint PDF $p(a, b, c)$
- ▶ assume we get $p(c), p(b|c), p(a|b, c)$ from domain expert
- ▶ define the joint PDF $p(a, b, c) := p(a|b, c)p(b|c)p(c)$
- ▶ define all partial joints by integration, e.g.

$$p(a, b) := \int p(a, b, c) dc$$

- ▶ now all sum rules hold! (also for $p(c)$ used in the definition)
- ▶ define the conditional PDF as quotients, e.g.

$$p(c|a) := p(a, c)/p(a)$$

- ▶ now all product rules hold (also for $p(a|b, c), p(b|c)$)

All rules hold by definition!

Random variable

Definition 5.2 离散随机变量 X

A *discrete random variable* X is a variable with values x ranging over some discrete set \mathcal{X} and a probability mass function (PMF)

$$f : \mathcal{X} \rightarrow [0, 1]$$

$$x \mapsto f(x)$$

概率质量函数 (PMF)

that sums up to one, $\sum_{x \in \mathcal{X}} f(x) = 1$.

Definition 5.3

A *continuous random variable* X is a variable with values x ranging over some continuous set \mathcal{X} and a probability density function (PDF)

$$f : \mathcal{X} \rightarrow \mathbb{R}_{\geq 0}$$

$$x \mapsto f(x)$$

that integrates to one, $\int_{x \in \mathcal{X}} f(x) dx = 1$.

RV用大写字母 X 表示，它的值用小写字母 x
更确切地说（也很复杂）：小写的 x 也是一个变量。

Note:

- ▶ RV are denoted by capital letters X , its values by small letters x
- ▶ More precisely (and complicated): the small x is a variable as well, that ranges over the values of the *random variable* X .
- ▶ Using the same letter, we know which small letter variable corresponds to which random variable (denoted by a capital letter).
- ▶ So, when we write $p(x)$ we are talking about the PDF with input x that belongs to random variable X . When we write $p(y) \dots$

Example – infer probability of wearing glasses

例子 - 推断戴眼镜的概率

Example — inferring probability of wearing glasses (1)

Question:

What's the probability that a person wears glasses?

一个人戴眼镜的概率是多少?

Example — inferring probability of wearing glasses (2)


Represent all unknowns as random variables (RVs) 将所有未知数表示为随机变量

- ▶ probability to wear glasses is represented by RV Y 戴眼镜的概率用RV Y 表示
- ▶ five observations are represented by RVs X_1, X_2, X_3, X_4, X_5 Possible values of the RVs 五个观察值用-----表示 RV 的可能值

▶ Y takes values $y \in [0, 1]$ (it is a probability value) Y 的取值 $y \in [0, 1]$ (它是一个概率值)。

- ▶ X_1, X_2, X_3, X_4, X_5 are binary, i.e. values 0 and 1

Graphical representation



Generative model and joint probability

似然反映的是：已知结果，反推原因

- ▶ we abbreviate $Y = y$ as y , $X_i = x_i$ as x_i
- ▶ $p(y)$ is the prior of Y , written fully $p(Y = y)$ $p(y)$ 是 Y 的先验，完全写成 $p(Y = y)$
- ▶ $p(x_i | y)$ is the likelihood of observation x_i $p(x_i | y)$ 是观察点 x_i 的似然性
- ▶ note that the likelihood is a function of y 注意，这个似然是 Y 的一个函数

Example — inferring probability of wearing glasses (3)

Probability of wearing glasses without observations

在没有观察的情况下，佩戴眼镜的概率

$$p(y|\text{"nothing"}) = p(y)$$

Probability of wearing glasses after one observation

观察一次后戴眼镜的概率

$$p(y|x_1) = Z_1^{-1} p(x_1|y)p(y)$$

Probability of wearing glasses after two observations

$$p(y|x_1, x_2) = Z_2^{-1} p(x_2|x_1, y)p(x_1|y)p(y) = Z_2^{-1} p(x_2|y)p(x_1|y)p(y)$$

...

Probability of wearing glasses after five observations

$$p(y|x_1, x_2, x_3, x_4, x_5) = Z_5^{-1} \left(\prod_{i=1}^5 p(x_i|y) \right) p(y)$$

Example — inferring probability of wearing glasses (4)

What is the likelihood?

$$p(x_1|y) = \begin{cases} y & \text{for } x_1 = 1 \\ 1 - y & \text{for } x_1 = 0 \end{cases}$$

More helpful RVs:

N表示观察值为1的数量

- ▶ RV N for the number of observations being 1 (with values n)
- ▶ RV M for the number of observations being 0 (with values m)

Probability of wearing glasses after five observations

$$\begin{aligned} p(y|x_1, x_2, x_3, x_4, x_5) &= Z_5^{-1} \left(\prod_{i=1}^5 p(x_i|y) \right) p(y) \\ &= Z_5^{-1} y^n (1 - y)^m p(y) \\ &= p(y|n, m) \end{aligned}$$

Example — inferring probability of wearing glasses (5)

Posterior after seeing five observations:在看到五次观察后的后方。

$$p(y|n, m) = Z_5^{-1} y^n (1 - y)^m p(y)$$

What prior $p(y)$ would make the calculations easy?

什么样的先验 $p(y)$ 会使计算变得容易?

$$p(y) = Z^{-1} y^{a-1} (1 - y)^{b-1} \quad \text{with parameters } a > 0, b > 0$$

called “Beta distribution with parameter a and b ” 称为“参数为 a 和 b 的贝塔分布”

Let's give the normalization factor Z of the beta distribution a name!

归一化相关系数

$$B(a, b) = \int_0^1 y^{a-1} (1 - y)^{b-1} dy$$

called “Beta function with parameters a and b ”

Note: for $a = 1, b = 1$, the Beta distribution is the uniform distribution $p(y) = 1$ for $y \in [0, 1]$, zero elsewhere. 对于 $a=1, b=1$, Beta分布是均匀分布 $p(y)=1$ 对于 $y[0, 1]$, 其他地方为零。

Gamma function, Beta function, and all that

from http://en.wikipedia.org/wiki/Gamma_function

and http://en.wikipedia.org/wiki/Beta_function

Gamma function (extension of factorial function)

$$\Gamma(z) = \int_0^{\infty} e^{-t} t^{z-1} dt \quad \text{for } z \in \mathbb{C}$$

$$\Gamma(n) = (n-1)! = n!/n \quad \text{for } n \in \mathbb{N}$$

Beta function (extension of ...?)

$$B(x, y) = \int_0^1 t^{x-1} (1-t)^{y-1} dt$$
$$= \frac{\Gamma(x)\Gamma(y)}{\Gamma(x+y)} \quad \text{for } x, y \in \mathbb{C} \text{ with } x + \bar{x}, y + \bar{y} > 0$$

$$B(m, n) = \frac{(m-1)! (n-1)!}{(m+n-1)!} \quad \text{for } m, n \in \mathbb{N}$$

$$= \binom{m+n}{n}^{-1} \frac{m+n}{mn} \quad \text{binomial coefficient}$$

Example — inferring probability of wearing glasses (6)

The prior of the probability with parameters a and b :

$$p(y) = \frac{y^{a-1}(1-y)^{b-1}}{B(a,b)}$$

The likelihood of the observations:

$$p(n, m | y) = y^n (1-y)^m$$

The posterior with the beta prior:

$$p(y|n, m) = \frac{y^{n+a-1}(1-y)^{m+b-1}}{B(a+n, b+m)}$$

Note:

- ▶ If the prior and the posterior have the same form (here: both are Beta distribution), we call prior the *conjugate* prior for likelihood. See https://en.wikipedia.org/wiki/Conjugate_prior.

共轭先验

Expectations 期望

Integration with densities

Weighted average/sum

$$\sum_i w_i x_i$$

with weights w_i

Weighted integral

$$\int w(x) x \, dx$$

with weights/density $w(x)$

Expectations of a random variable (1)

Definition 5.4 (expected value of x)

1. The *expected value* of a continuous random variable X with PDF $p(x)$ is:

$$E X = E(X) = \int x p(x) dx$$

2. The *expected value* of a discrete random variable X with probability mass function $p(x)$ is:

$$E X = E(X) = \sum_x x p(x)$$

- ▶ The E turns a random variable X into a single number (usually into a single value from its range).

Expectations of a random variable (2)

Definition 5.5 (expected value of a function of x)

1. The *expected value* of a function f wrt the continuous random variable X with PDF $p(x)$ is:

$$E f(X) = E(f(X)) = \int f(x)p(x)dx$$

2. The *expected value* of a function f wrt the discrete random variable X with probability mass function $p(x)$ is:

$$E f(X) = E(f(X)) = \sum_x f(x)p(x)$$

- ▶ $f(X)$ can be also understood as a new random variable $Y = f(X)$.
- ▶ Sometimes we write E_x to specify which variable is summed up. Note that we use a small x in those cases since the subindex to E binds the variable and makes it local, e.g. $E_x x$.
- ▶ “wrt” = “with respect to”

Expectations of a random variable (3)

Examples

- ▶ Mean

$$E X = E(X) = E_x x$$

- ▶ Variance

$$\text{Var } X = \text{Var}(X) = E(X - \mu)^2 = E_x (x - \mu)^2$$

is the average squared distance to the mean $\mu = E_x x$

- ▶ Product- and sum-rule combined

$$p(y) = \int p(x, y) dx = \int p(y|x) p(x) dx = E_x p(y|x) = E p(y|X)$$

- ▶ Probabilities as expectations

$$p(X \in A) = E(1_A(X)) = E[X \in A]$$

with Iverson bracket $[F] = 1$ if formula F is true, and zero otherwise.